



21.03.19 12:55

Accelerating Machine Learning Inference with Probabilistic Predicates



Jetic Gū
CMPT 843

Overview

- Focus: Query Optimisation
- Architecture: Relational Data Platforms with UDFs
- Core Ideas:
 1. Observation: expensive UDFs can seriously influence performance
 2. Propose: probabilistic predicates using fast ML methods to filter data before feeding into the UDF layer

Queries with UDF

- UDF: user defined functions
 - Example: ML function for object detection
 - Input: Video data/Image data in the database
 - Output: Object type (cheese), Object Attributes (colour=Yellow)
- Using UDF in queries are more frequent nowadays

Queries with UDF

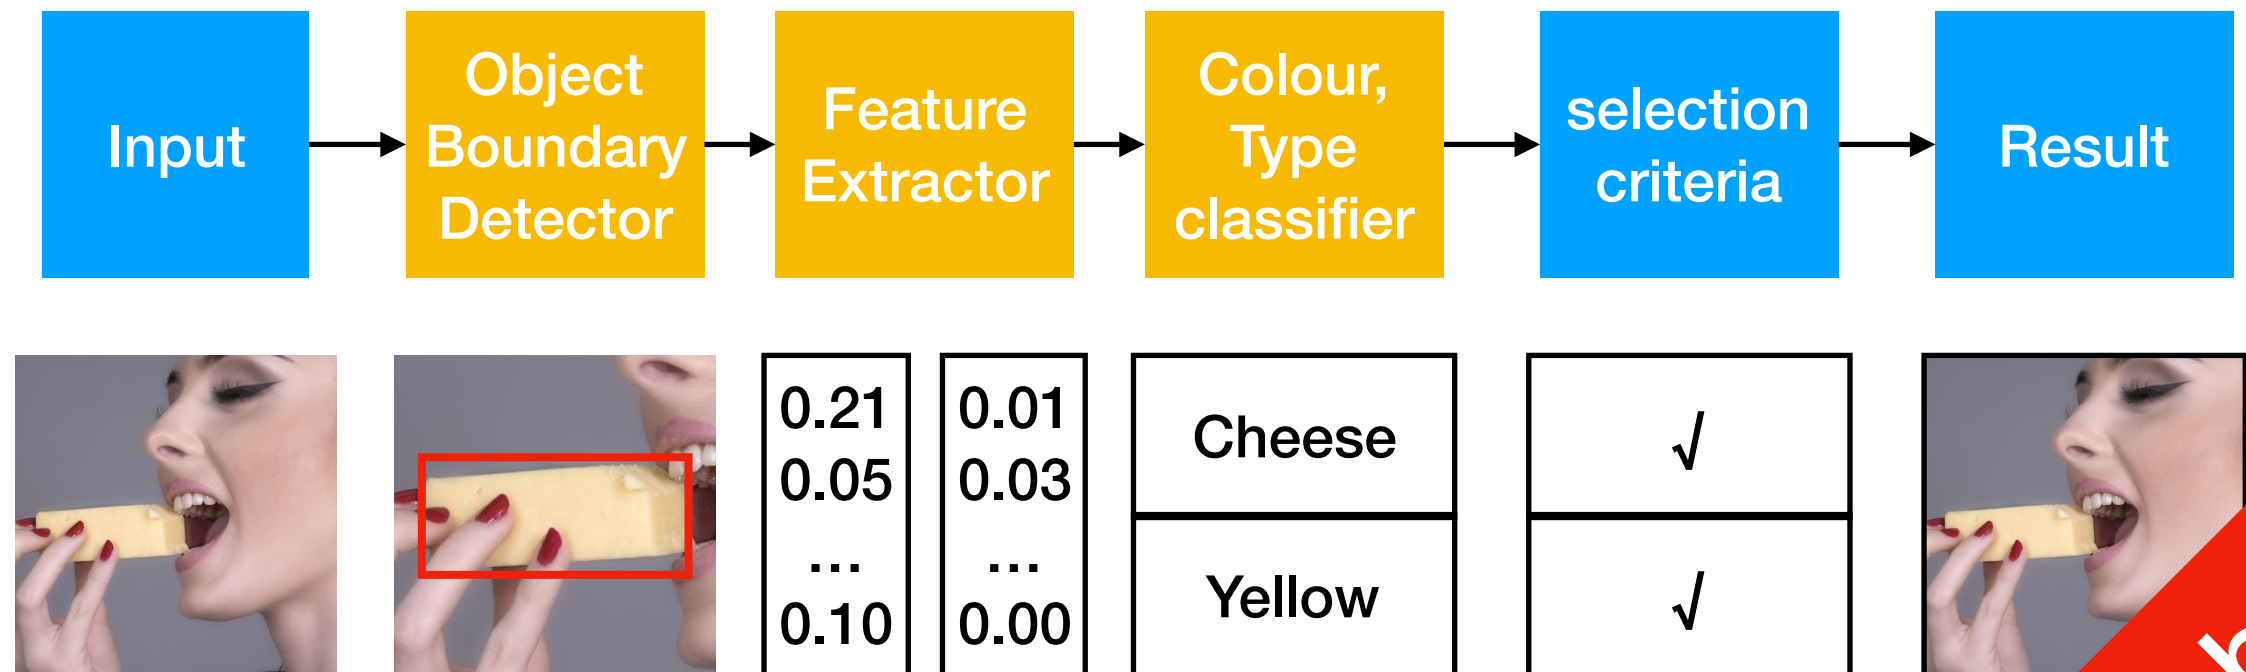
- Objective: get something within the DB
- Pipeline:



Example

Queries with UDF


- UDFs are more general. Multiple feature extraction steps and inference stages will be needed in practice
- Objective: return pictures with Yellow Cheese
- Pipeline:



Problem

Queries with UDF

- **Sloooooooooooooow** execution of UDFs
- Why? Because this entire chain of UDFs are run on **ALL** entries
- Traditional query optimisation techniques such as predicate pushdown are **NOT** useful here



Problem

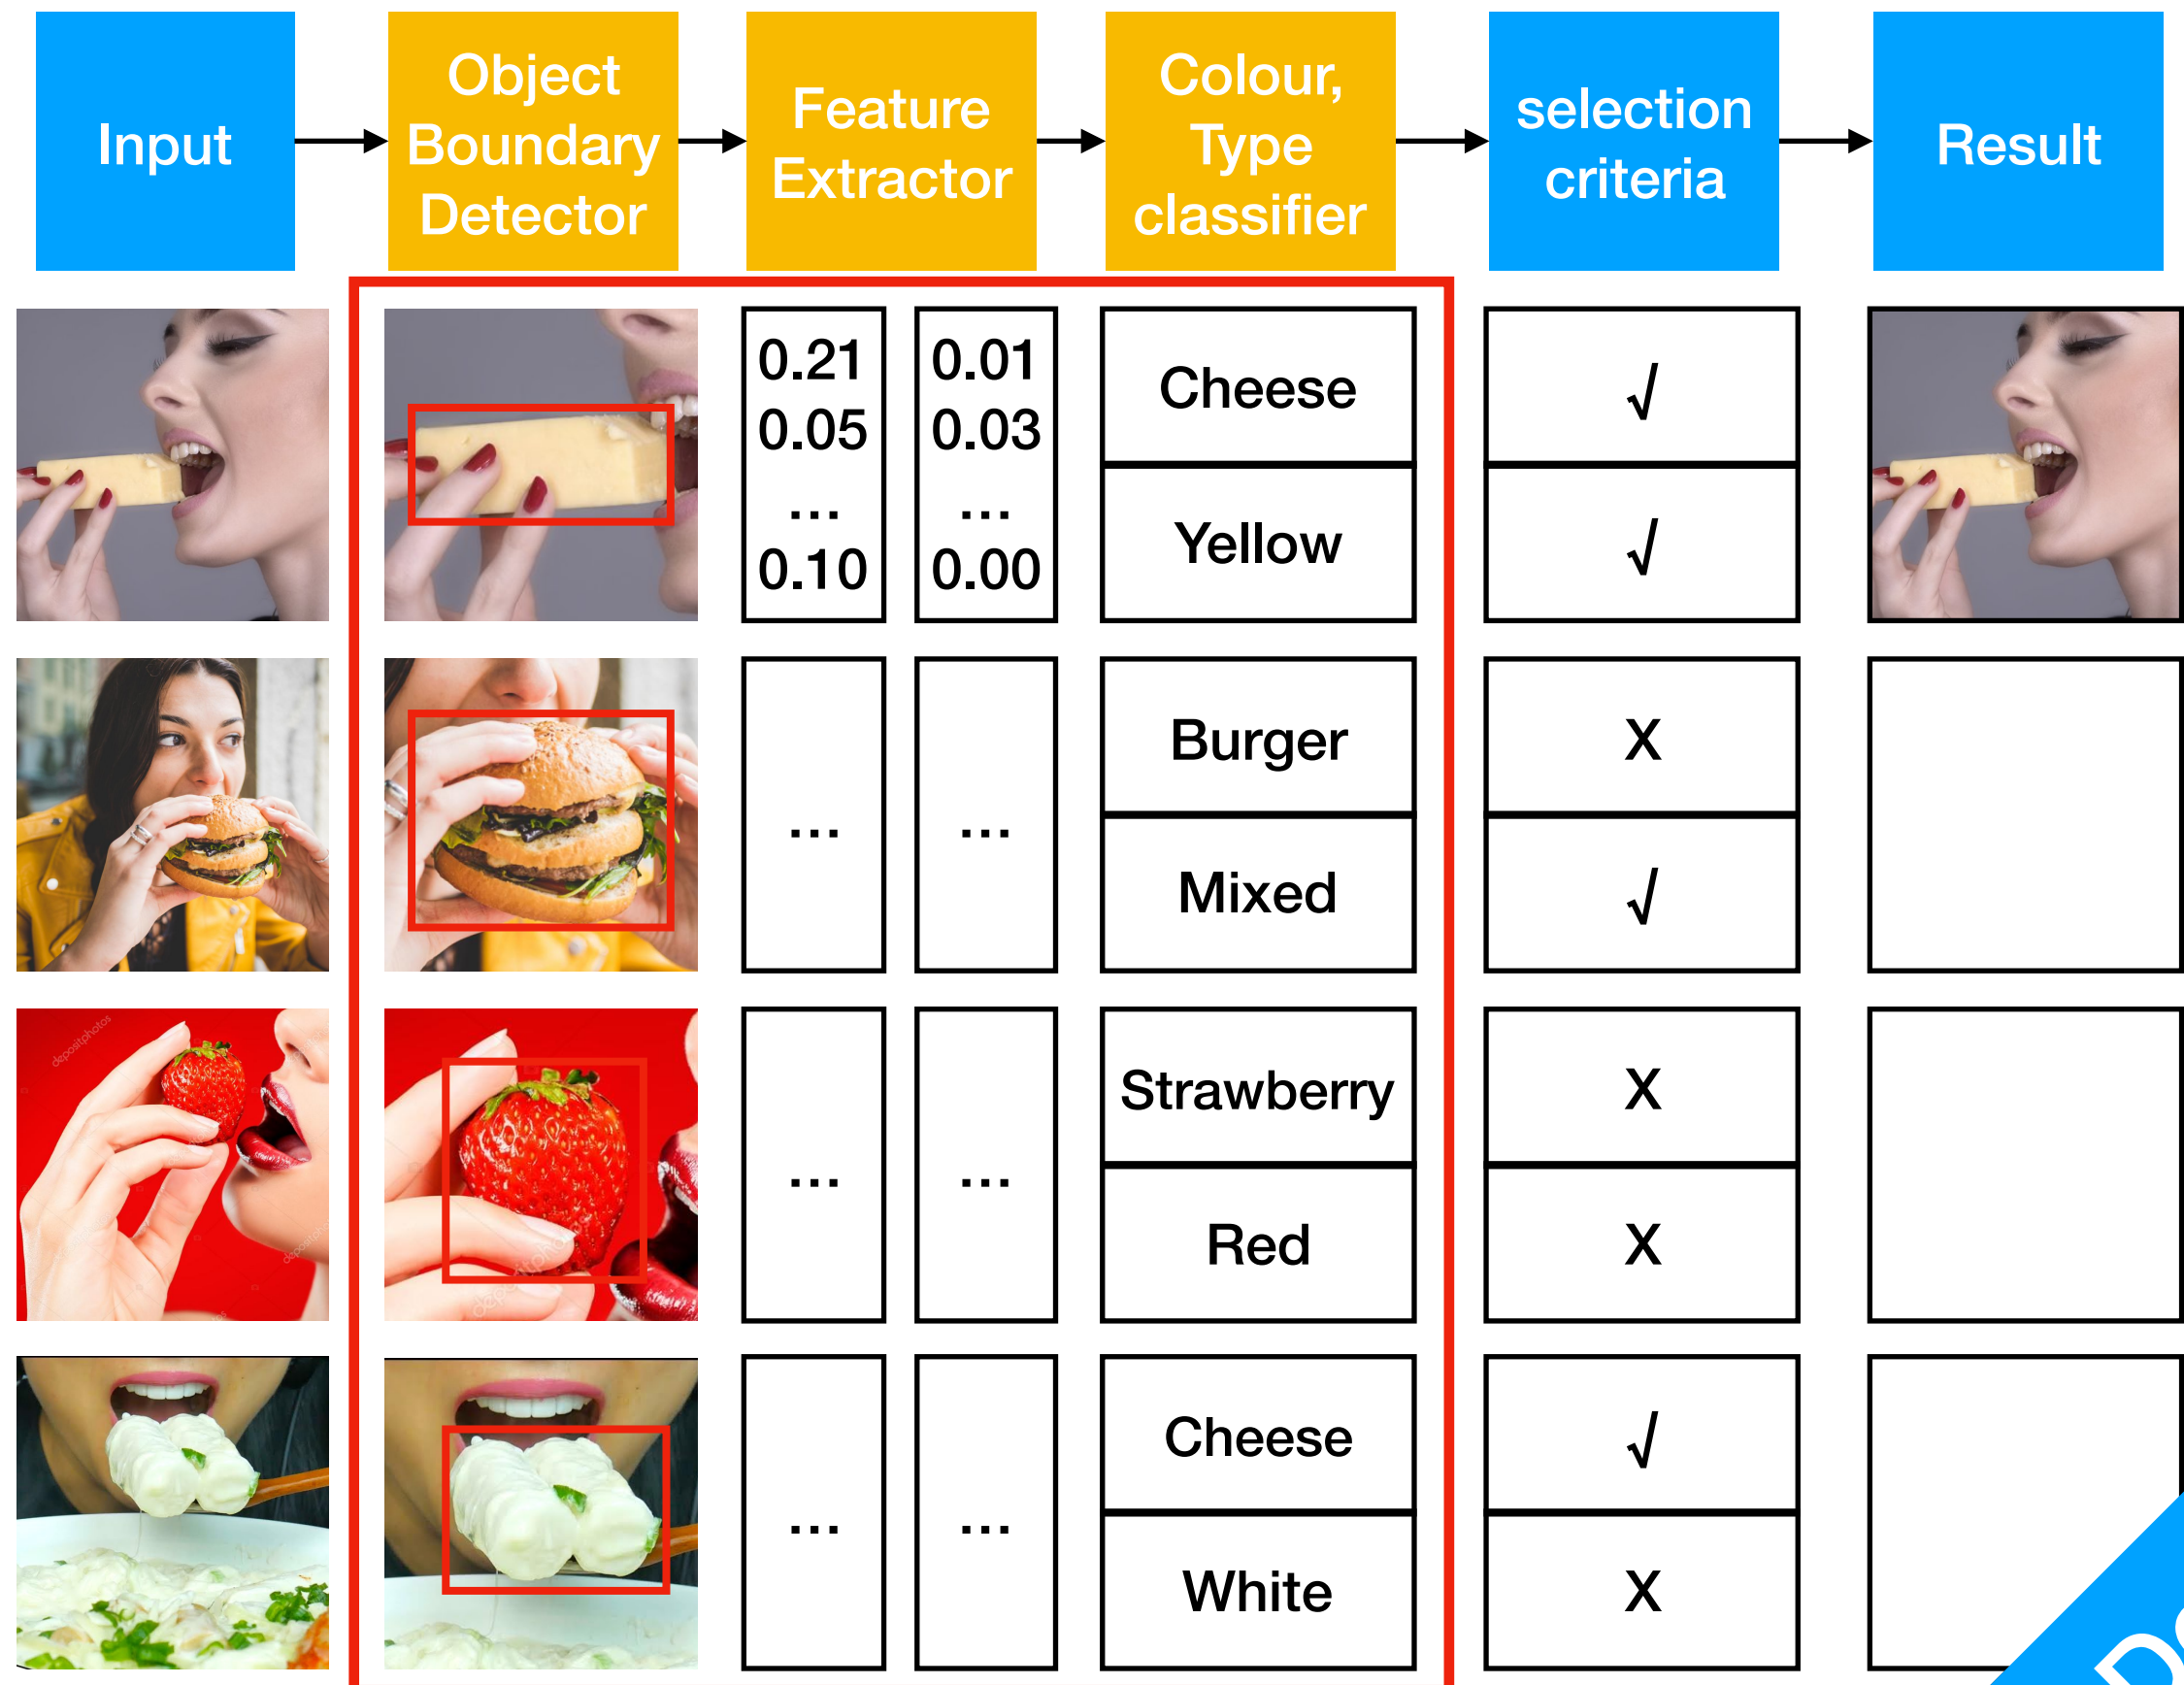
Probabilistic Predicates

- PPs are binary classifiers
- Get rid of highly unlikely entries before running through the UDFs.
- PPs are much more lightweight and error tolerant: high precision on negatives, low recall is OK.



Concept

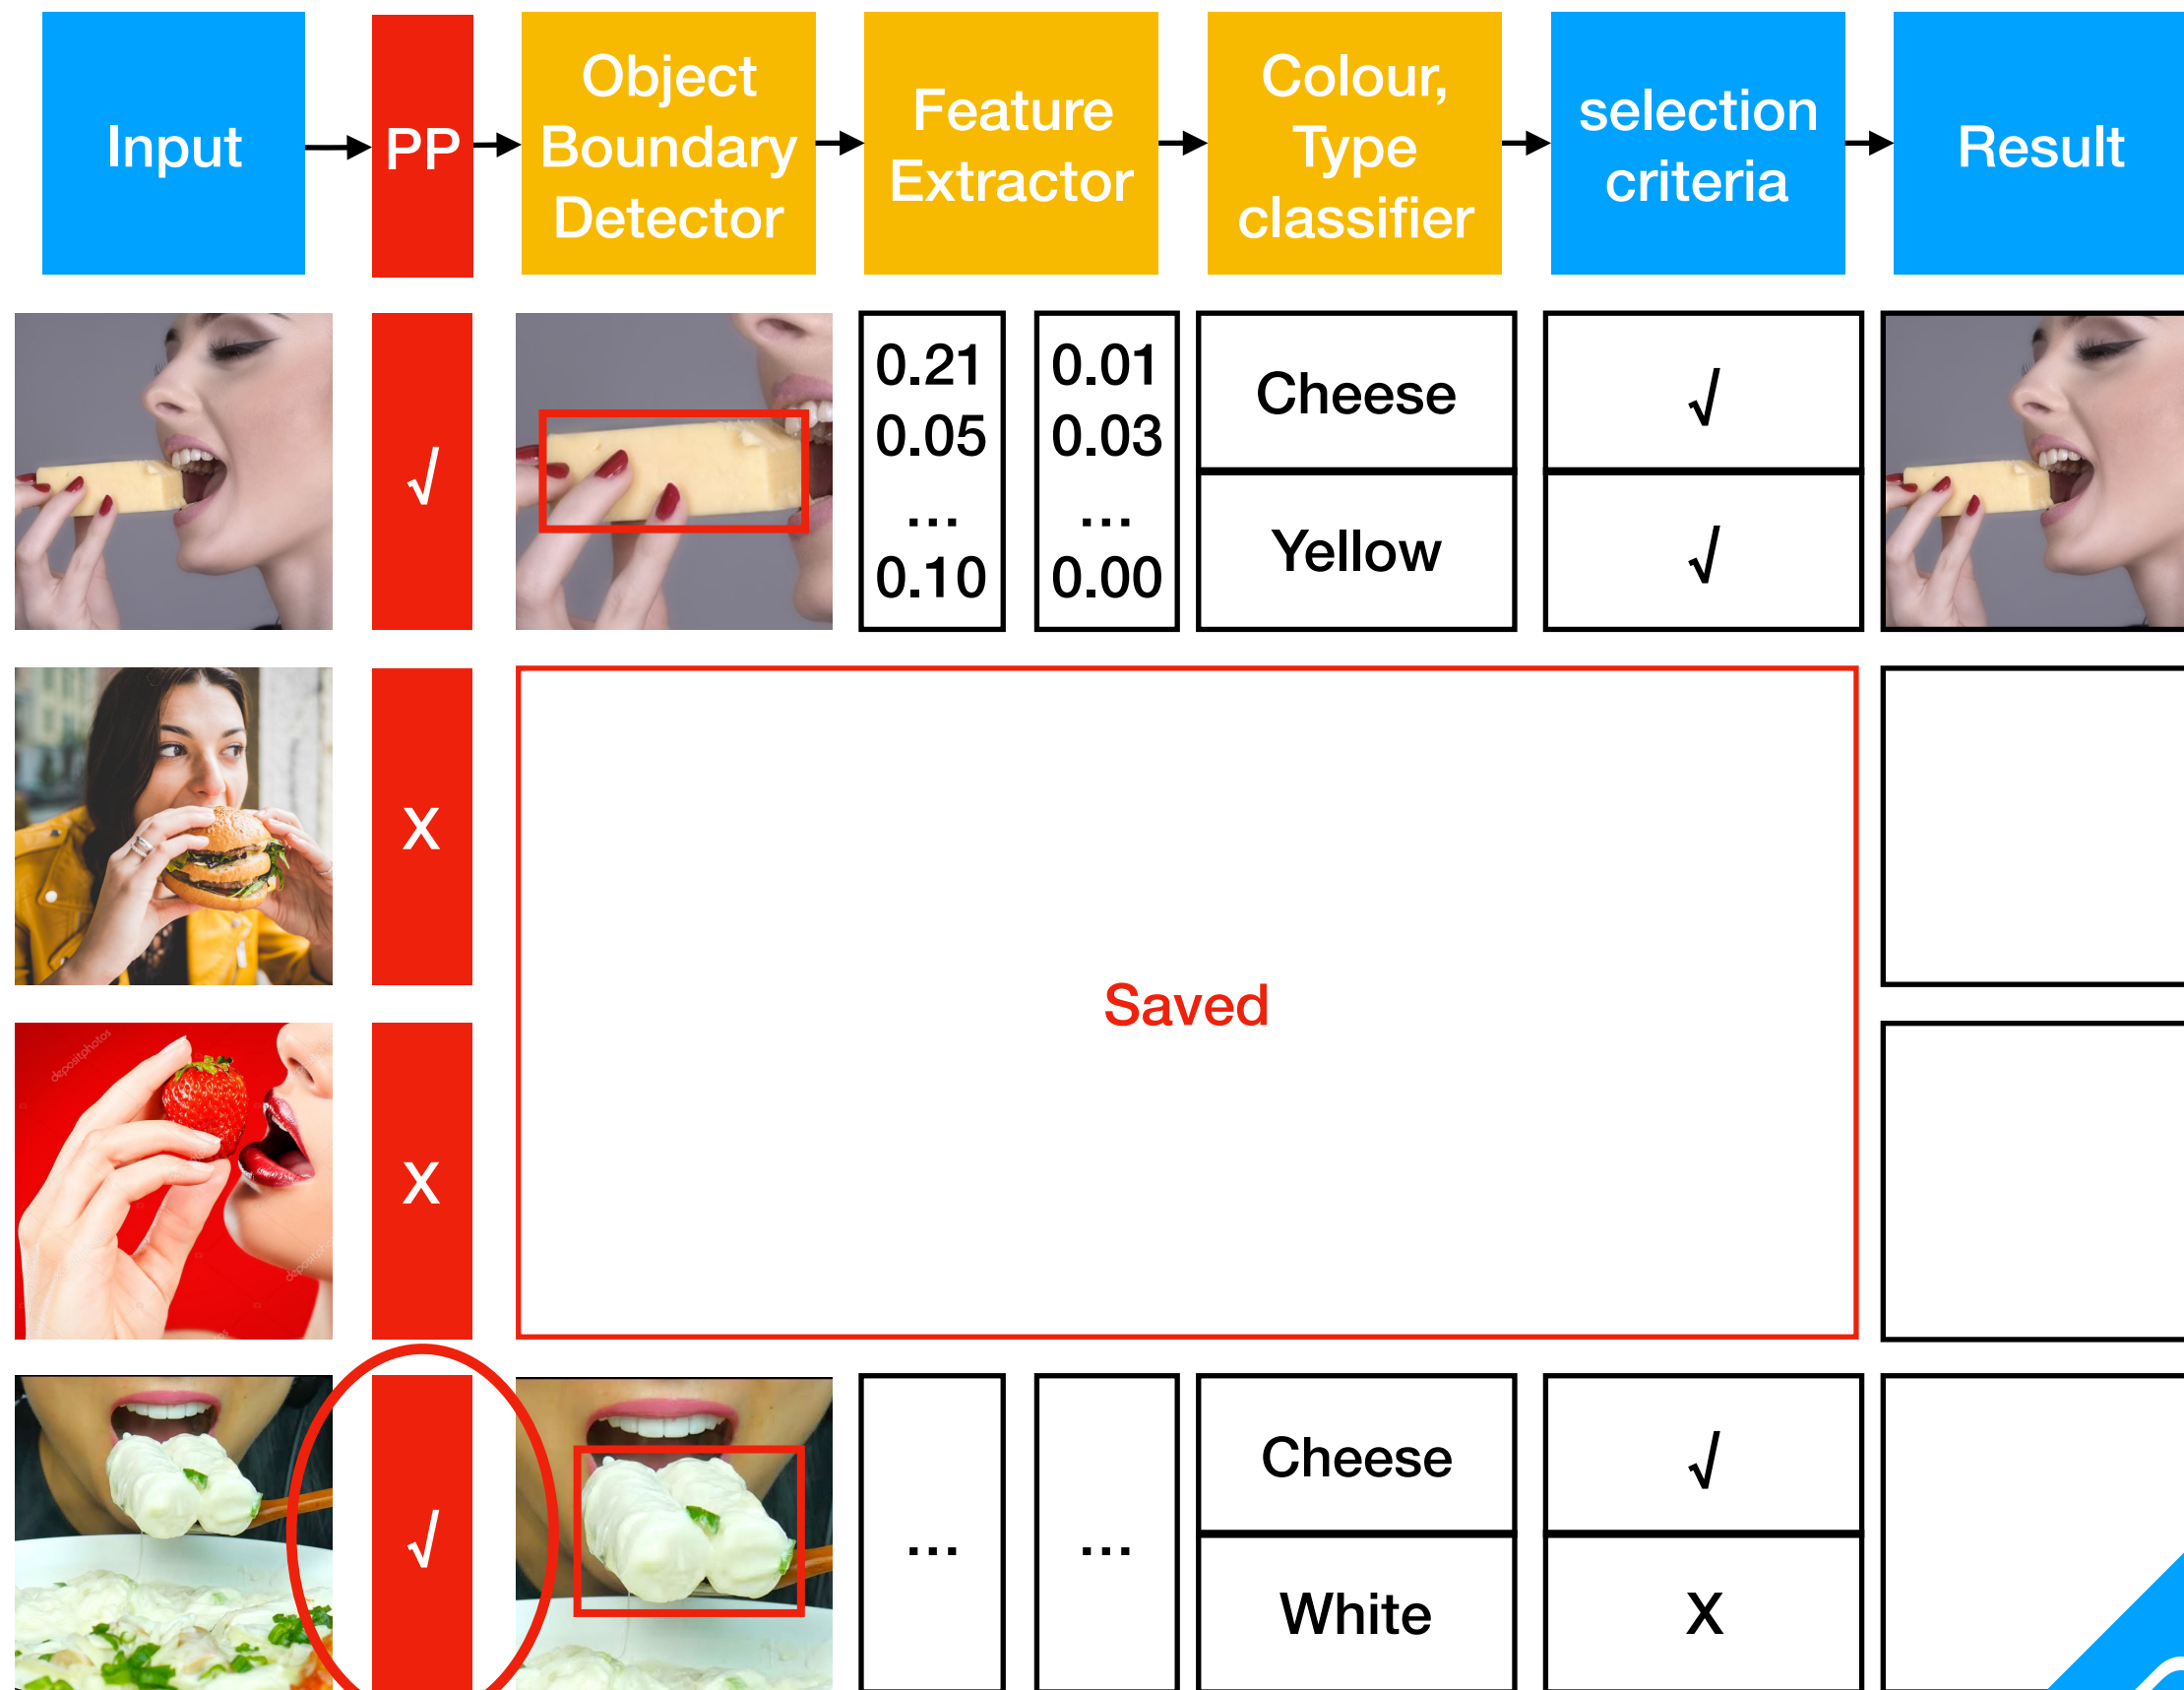
No Probabilistic Predicates



Query Optimiser cannot help, super slow execution chain

Demo

Probabilistic Predicates



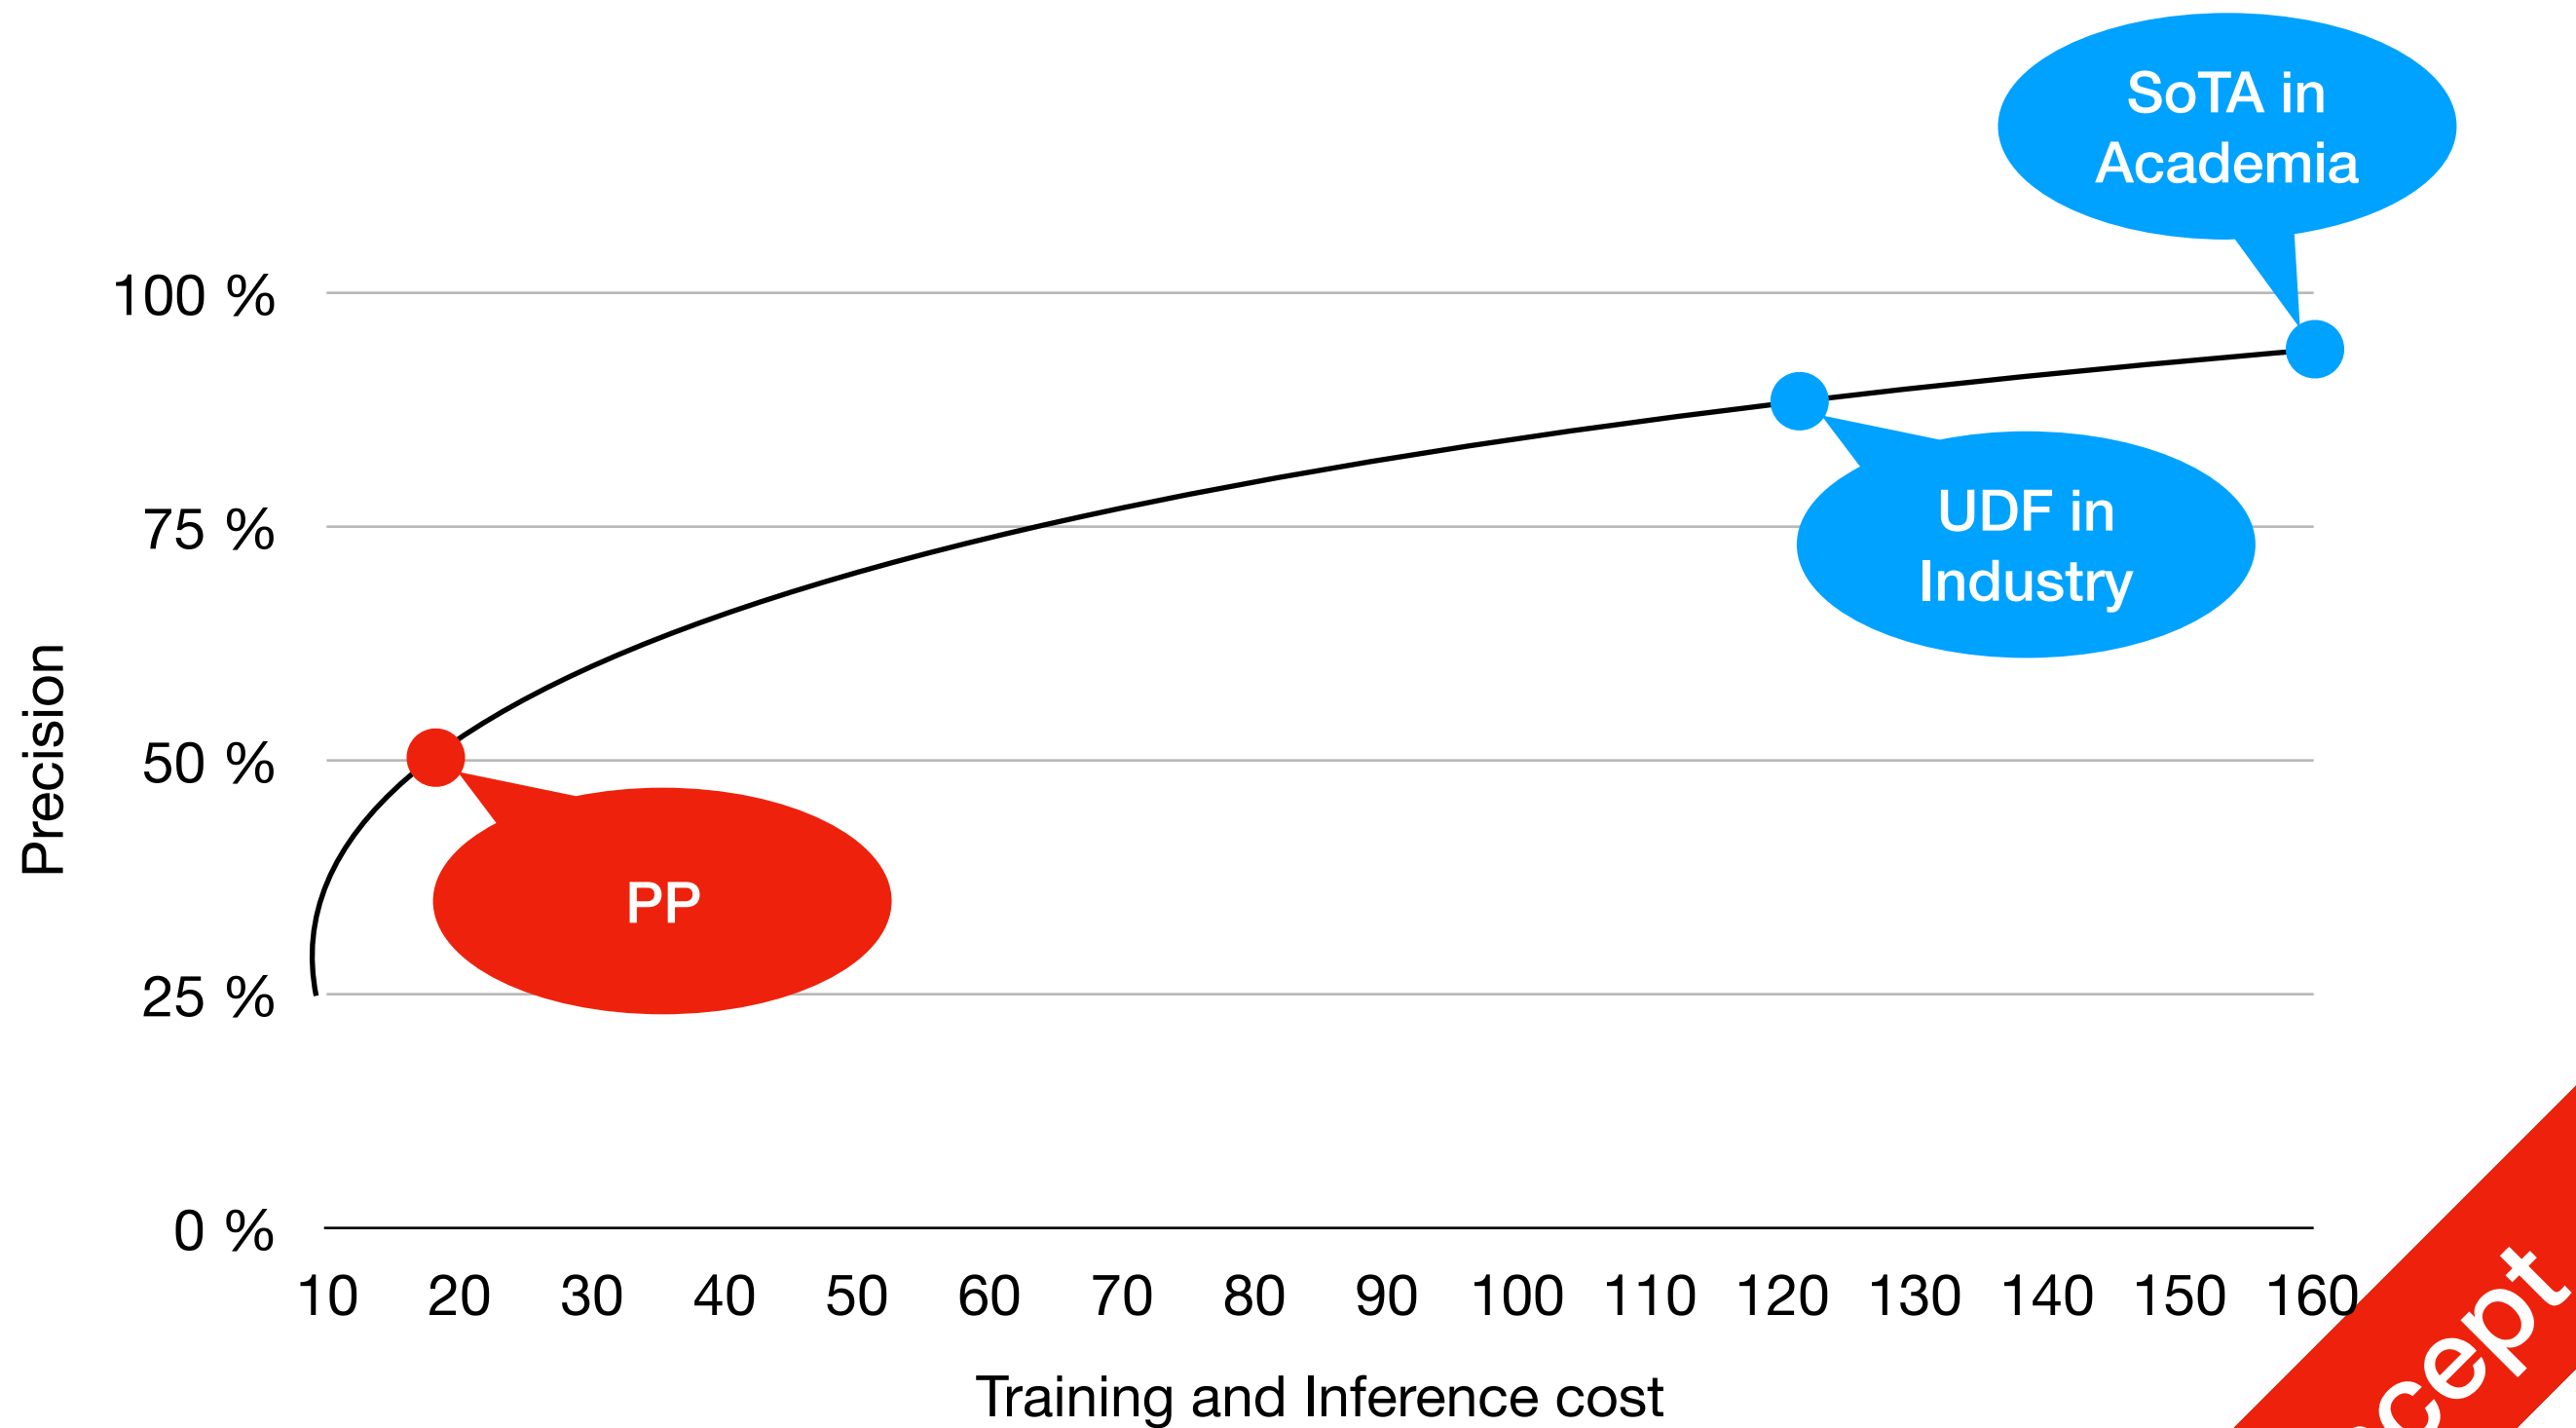
Incorrect, but that's OK

Demo

Comparison

- UDF
 - High execution cost
 - Long chain of execution
 - High precision
 - Give highly precise results
 - Training is costly
- PP
 - Very low execution cost
 - One time fast execution
 - High negative precision
 - Quickly eliminate highly unlikely entries
 - Cheap Training

Comparison



Concept

Comparison

	System	Features	Classifiers/Regressors	Materialization Cost (sec)	Query predicate	Selectivity
Online	Ads recommendations [42]	Bag-of-words	Collaborative Regressor	$10^{-2} - 10^{-1}$	1 binary	1-in-hundreds
	Video recommendations [16]	Browse history	Bayesian Regressor	$10^{-1} - 10^1$	1 binary	1-in-thousands
	Credit card fraud [47]	Physical loc. etc.	Neural Network	$10^{-2} - 10^{-1}$	1 binary	1-in-thousands
Offline	Video tagging [24]	Keypoints	SVM w/ RBF kernel	$10^{-1} - 10^1$	n categorical	1-in-thousands
	Spam filtering [6]	Bag-of-word	Naive Bayes Classifier	$10^{-2} - 10^{-1}$	1 binary	1-in-several
	Image tagging [37, 55]	Keypoints	Collaborative Regressor	$10^{-1} - 10^1$	n categorical	1-in-thousands

- PP runs 10 times faster than UDF chain, filtering out 60% of entries
- That's 2x speed up!

Let's get technical

Technical

Different PPs

- Computational Models, use model selection techniques at inference time
 - linear support vector machine
 - kernel density estimator
 - deep neural networks
- Define and train one PP per simple clause $PP_p = D, m, r[a]$
 - Each PP has a reduction rate $r[a]$, where target accuracy is a .
 - Pre-sampling different a values, achieved using training data D . m indicates the computational model type.
 - Dimension reduction techniques are used to speed up

Reduction Rate and Target Accuracy

- The output of a PP is a probability distribution:

$$P(\mathbf{rejected} \mid e, p) \in [0,1]$$

- Essentially, how likely the PP thinks this entry e will not satisfy the predicate p
- Activation threshold t : if $P(\mathbf{reject}) \geq t$, filter it out
- By adjusting the activation threshold (how much we want to trust the PP to throw things out), we can get $r[a]$

Choosing Simple Clauses to Define PP

- Inferred from historical queries
- Define and train one PP per simple clause
 - Use query optimiser to assemble PPs

Assembly (Query Optimiser)

- Convert complex query expression to logical expressions using PPs (not strict it's just filtering here)

Complex predicate	Implied logical expr. over PPs
$(p \vee q) \wedge \neg r \wedge \mathcal{P}_{\text{rem}}$	$\Rightarrow p \vee q \Rightarrow \text{PP}_{p \vee q} \Rightarrow \text{PP}_p \vee \text{PP}_q$ $\Rightarrow \neg r \Rightarrow \text{PP}_{\neg r}$ $\Rightarrow \text{PP}_{(p \vee q) \wedge \neg r} \Rightarrow (\text{PP}_p \vee \text{PP}_q) \wedge \text{PP}_{\neg r}$ $\Rightarrow \text{PP}_{(p \wedge \neg r) \vee (q \wedge \neg r)} \Rightarrow \text{PP}_{p \wedge \neg r} \vee \text{PP}_{q \wedge \neg r} \Rightarrow$ $(\text{PP}_p \wedge \text{PP}_{\neg r}) \vee (\text{PP}_q \wedge \text{PP}_{\neg r})$

omitted, and that's OK

- Use query optimiser to
 - Determine for each PP, which CM m to use, within what target accuracy range a to reduce overall $r[a]$
 - Determine the execution order

Assembly (Query Optimiser)

$$(PP_p \wedge PP_{\neg r}) \vee (PP_q \wedge PP_{\neg r})$$

- 3 subproblems.

1. different allocations of the query's accuracy budget to individual PPs

$$PP_p(r) = \begin{bmatrix} (a = 100\%, r = 0, m = \mathbf{None}) \\ (a = 95\%, r = 15\%, m = \mathbf{rule}_p) \\ (a = 90\%, r = 20\%, m = \mathbf{SVM}_p) \\ (a = 85\%, r = 40\%, m = \mathbf{NN}_p) \end{bmatrix} \quad PP_{\neg r}(r) = \begin{bmatrix} (a = 100\%, r = 0, m = \mathbf{None}) \\ (a = 95\%, r = 10\%, m = \mathbf{rule}_p) \\ (a = 90\%, r = 15\%, m = \mathbf{SVM}_p) \\ (a = 85\%, r = 50\%, m = \mathbf{NN}_p) \end{bmatrix}$$

Given accuracy target, use dynamic programming to maximise reduction rate.

$$PP_{(PP_p \wedge PP_{\neg r}) \vee (PP_q \wedge PP_{\neg r})}(r)[a \geq A] = ?$$

Demo

Assembly (Query Optimiser)

$$(PP_p \wedge PP_{\neg r}) \vee (PP_q \wedge PP_{\neg r})$$

- 3 subproblems.
 1. different allocations of the query's accuracy budget to individual PPs
 2. different orderings of PPs within a conjunction or disjunction

to reduce the execution cost here

Assembly (Query Optimiser)

$$(PP_p \wedge PP_{\neg r}) \vee (PP_q \wedge PP_{\neg r})$$

- 3 subproblems.
 1. different allocations of the query's accuracy budget to individual PPs
 2. different orderings of PPs within a conjunction or disjunction
 3. compute the cost and reduction rate of the resulting plan for query optimiser to work on

PP Experiments

Dataset	Approach	Avg. data reduction \bar{r} for accuracy a		
		$\bar{r}[1]$	$\bar{r}[0.99]$	$\bar{r}[0.9]$
UCF101	PCA+KDE	0.47	0.56	0.64
	PCA + SVM	0.35	0.45	0.54
	Raw + SVM	0.35	0.47	0.59
COCO	DNN	0.28	0.50	0.83
	SVM		0.31	
ImageNet	DNN	0.71	0.84	0.96
	SVM		0.39	
	DNN trained on COCO	0.25	0.49	0.82


- Individual PP evaluation (recognising objects on images)

PP Experiments

Dataset	Approach	PP cost to ...		Optimality for a	
		Train (per 1K rows)	Test	$a = 1$	$a = 0.9$
UCF101	PCA+KDE	14s	3ms	0.55	0.77
LSHTC	FH + SVM	1s	1ms	0.29	0.87
COCO	DNN*	110s	10ms	0.28	0.83

- Individual PP evaluation (recognising objects on images)
- The latency to train and test PPs of different types

End2End Experiments



- Evaluating TRAF20 query set on 100 GB online data

End2End Experiments

ID	PP cons.	#PPs	PP inf.	Sub.UDF	Selectivity	Reduction
4	27s	1	2ms	23ms	0.67	11%
8	68s	2	5ms	55ms	0.41	20%
20	155s	4	12ms	85ms	0.01	60%
Avg.	79s	2.5	6ms	52ms	0.20	59%

- Evaluating TRAF20 query set on 100 GB online data
- Training and inference overhead for deploying PPs in online machine learning query processing.
- ID indicates the individual query they used for this analysis

- Most query optimisation: optimally ordering the existing predicates in the query
- When predicates rely on columns generated by user- defined operators, shows that performance-optimal ordering of the UDFs and predicates is NP-hard.
- Approximate predicates use the same relation as the query predicates and are not for blobs
- PP
 - trained without any knowledge of the inference modules that are used in a query
 - can be applied on non-relational datasets

Thank you.